



6BUIS001W  
Business  
Intelligence –  
Coursework 1

Vasileios Protopapas-  
W1735144



## Contents

1 <sup>st</sup> Objective (partitioning clustering) .....	1
Discussion of methodologies used for reducing the input dimensionality .....	1
Pre-processing tasks .....	1
Scaling .....	1
Outlier Removal .....	2
Defining all cluster centres .....	4
K-Means Analysis & final “winner” cluster case with providing brief explanation .....	9
Coordinates for each centre clustering group .....	12
2 <sup>nd</sup> Objective (MLP) .....	13
References .....	13
Appendix .....	13

## 1<sup>st</sup> Objective (partitioning clustering)

### Discussion of methodologies used for reducing the input dimensionality

One of the methodologies that are normally used to reduce input dimensionality are aspects such as missing values ratio. This is responsible for finding any data value that are missing from their data columns. As mentioned by Filter, L.V and Filter in regards to the missing values ratio “The goal, then, becomes to remove those data columns with too many missing values”. Another method is the low variance filter which like its counterpart of MVR it’s purpose is to also remove data from columns whose variance is too low, however this requires normalization. Another method would be the high correlation filter which take in values from data columns that are holding similar trends with specific pairs that hold this attribute being reduced to one. Principal Component Analysis is an important methodology as it purpose is to project all the data points of a column and only point it to a few specific components to reduce dimensionality but also ensure that no data variation is lost in the process also mentioned by (Karamizadeh, S., Abdullah, S.M., Manaf, A.A., Zamani, M. and Hooman, A., 2013.) on the overview on PCA they mentioned that its used to “convert a set of observations of possibly correlated variables into a set of values of linearly” helping to further show it’s importance to lower input dimensionality.

### Pre-processing tasks

#### Scaling

#Scaling and centralizing the data

```
winedata <- winedata
```

#Removes the 12th variable

```
winedata1 = winedata[,-12]
```

winedata1		4873 obs. of 11 variables			
	fixed.acidity	volatile.acidity	citric.acid	residual.sugar	
1	6.2	0.450	0.26	4.4	
2	9.8	0.360	0.46	10.5	
3	5.5	0.485	0.00	1.5	
4	6.4	0.595	0.14	5.2	
5	7.6	0.480	0.37	0.8	
6	7.2	0.320	0.47	5.1	
	chlorides	free.sulfur.dioxide	total.sulfur.dioxide	density	
1	0.063	63	206	0.9940	
2	0.038	4	83	0.9956	
3	0.065	8	103	0.9940	
4	0.058	15	97	0.9951	
5	0.037	4	100	0.9902	
6	0.044	19	65	0.9910	
	pH	sulphates	alcohol		
1	3.27	0.52	9.8		
2	2.89	0.30	10.1		
3	3.63	0.40	9.7		
4	3.38	0.36	9.0		
5	3.03	0.39	11.4		
6	3.03	0.41	12.6		

It is shown that the quality variable is no longer there as a result.

#### Outlier Removal

# checking for missing values

```
sum(is.na(winedata1))
```

```
summary(is.na(winedata1))
```

```
> summary(is.na(winedata1))
fixed.acidity    volatile.acidity    citric.acid
Mode :logical    Mode :logical    Mode :logical
FALSE:4873       FALSE:4873       FALSE:4873
residual.sugar   chlorides         free.sulfur.dioxide
Mode :logical    Mode :logical    Mode :logical
FALSE:4873       FALSE:4873       FALSE:4873
total.sulfur.dioxide density          pH
Mode :logical    Mode :logical    Mode :logical
FALSE:4873       FALSE:4873       FALSE:4873
sulphates        alcohol
Mode :logical    Mode :logical
FALSE:4873       FALSE:4873
```

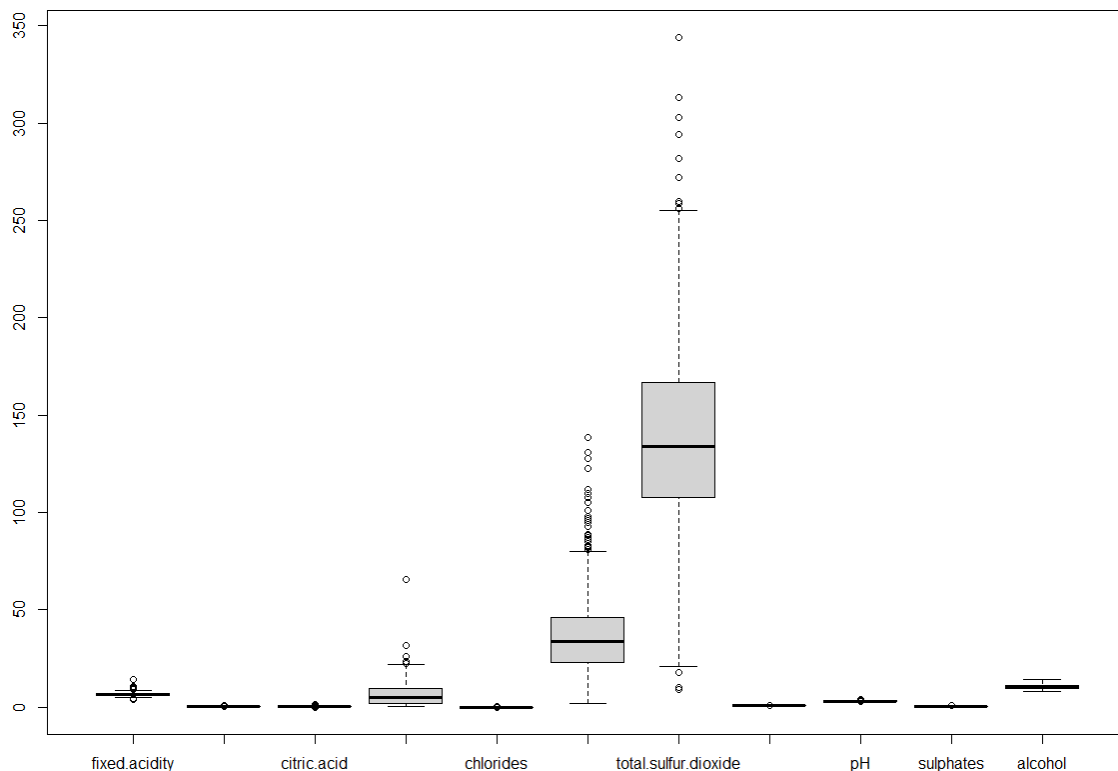
Here I have checked to see if there are any missing values with the “sum(is.na(data))” function. As shown in the image, this dataset contains no missing values

#Outlier detection by storing all numerical variables in an array structure

```
col = c("fixed.acidity","volatile.acidity","citric.acid", "residual.sugar","chlorides","free.sulfur.dioxide",  
,"total.sulfur.dioxide","density","pH","sulphates", "alcohol")
```

#Using the BoxPlot function to detect any presence of outliers in the data columns.

```
boxplot(winedata1[,c("fixed.acidity","volatile.acidity","citric.acid"  
,"residual.sugar","chlorides","free.sulfur.dioxide", "total.sulfur.dioxide","density","pH","sulphates",  
"alcohol")))
```



From the visuals created by the boxplot this shows that all the variables contain outliers in the data values apart from the 'alcohol' variable

```
> #Checking whether the outliers have been replaced by null d  
a or not  
> as.data.frame(colsums(is.na(winedata1)))  
               colsums(is.na(winedata1))  
fixed.acidity                113  
volatile.acidity             182  
citric.acid                  270  
residual.sugar                 7  
chlorides                    206  
free.sulfur.dioxide           46  
total.sulfur.dioxide           15  
density                        5  
pH                             75  
sulphates                    124  
alcohol                       0
```

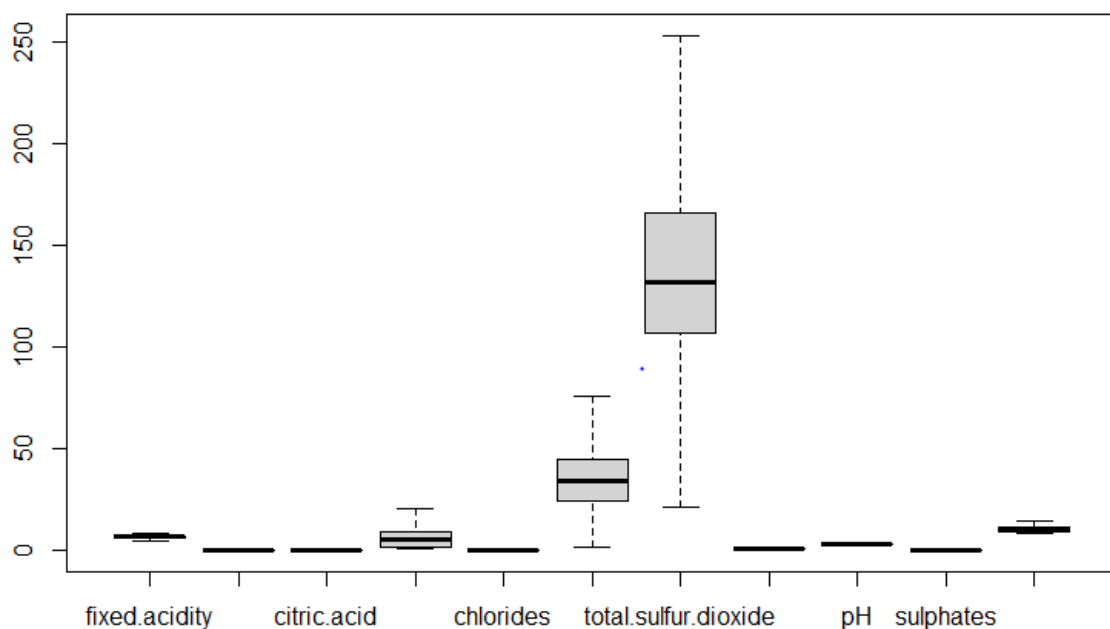
Here we can see that all the variables apart from alcohol that have outlier data

```

> #This is responsible for removing all the null values
> winedata1 = drop_na(winedata1)
> as.data.frame(colSums(is.na(winedata1)))
      colSums(is.na(winedata1))
fixed.acidity                0
volatile.acidity             0
citric.acid                  0
residual.sugar               0
chlorides                    0
free.sulfur.dioxide          0
total.sulfur.dioxide         0
density                      0
pH                           0
sulphates                    0
alcohol                      0

```

Here all the data has been turned to NULL and has been removed



After this operation all the outliers present on the first boxplot are now all gone

### Defining all cluster centres

#Manually finding clusters 1

```

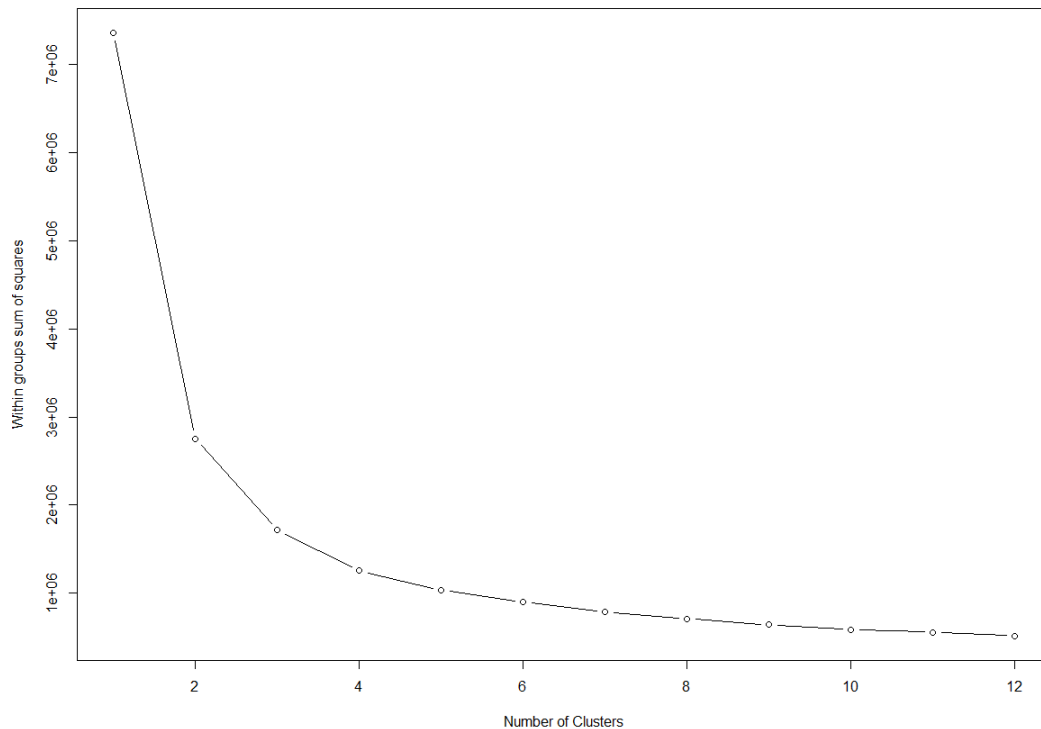
wss <- 0
for (i in 1:12){
  wss[i] <-
    sum(kmeans(winedata1, centers=i)$withinss)
}
plot(1:12,
     wss,

```

```

type="b", ### "b" for both####
xlab="Number of Clusters",
ylab="Within groups sum of squares")

```



This graph shows that k=3 is the best possible cluster for the data

```

#Manually finding clusters in a larger criteria
set.seed(1234)

```

```

nc <- NbClust(winedata1,
              min.nc=2, max.nc=15,
              method="kmeans")

```

```

table(nc$Best.n[1,])

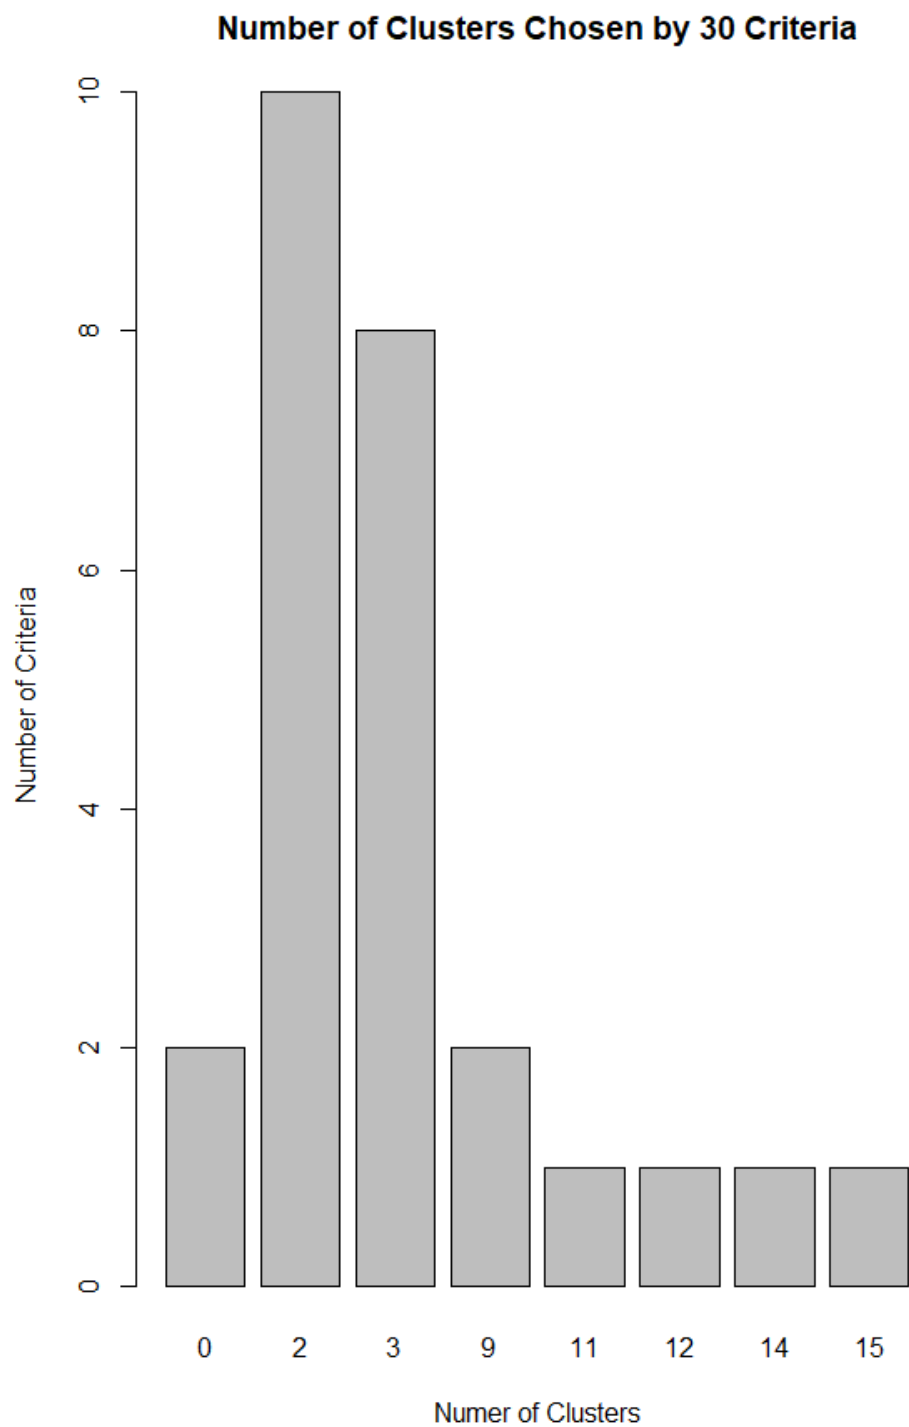
```

```

barplot(table(nc$Best.n[1,]), # provide bar charts####
        xlab="Nuner of Clusters",

```

```
ylab="Number of Criteria",  
main="Number of Clusters Chosen by 30 Criteria")
```

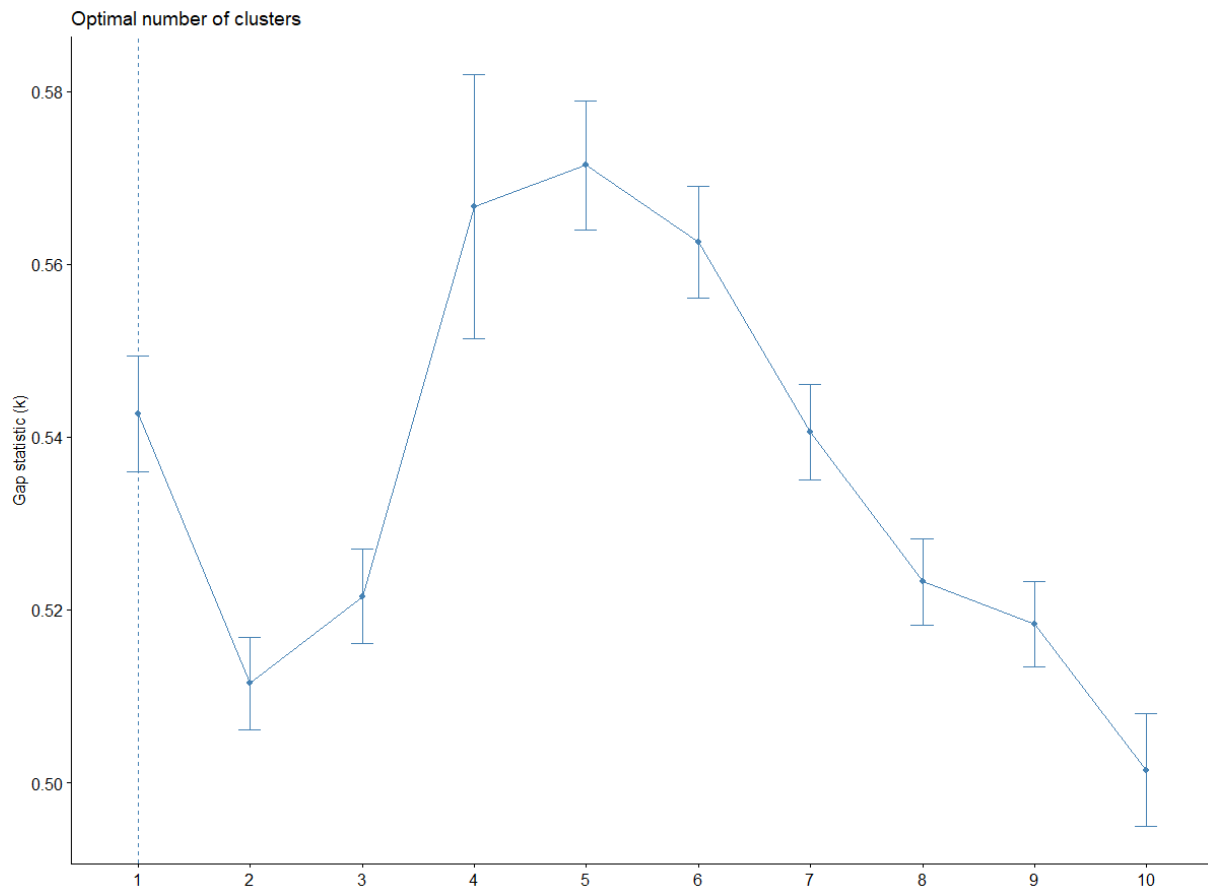


Here we can see that  $K=2$  is the best number

Here is the dominant winner for the 2 clusters

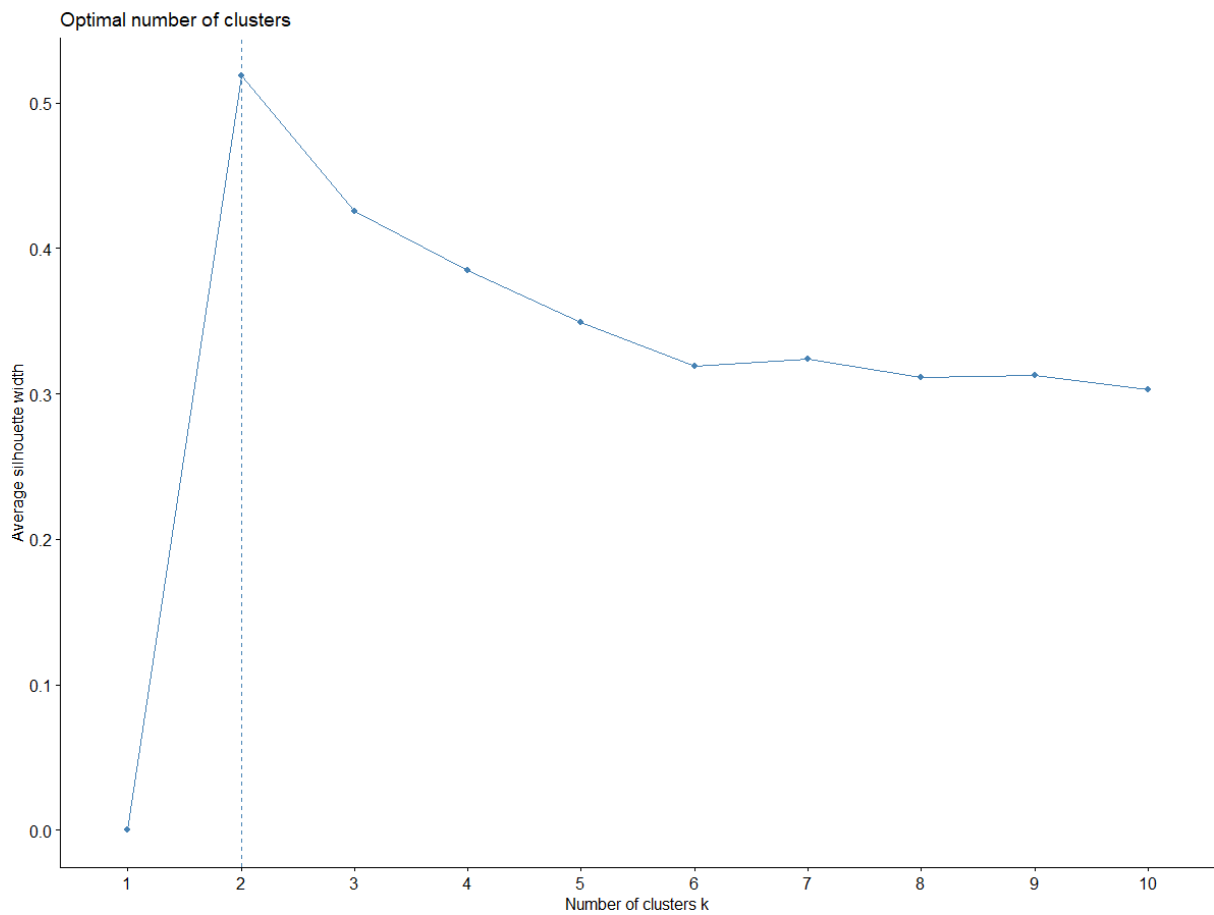
# Automatically finding clusters from 3 different methods

```
fviz_nbclust(winedata1[,ncol(winedata1)], FUNcluster = kmeans, method = 'wss')  
fviz_nbclust(winedata1[,ncol(winedata1)], FUNcluster = kmeans, method = 'silhouette')  
fviz_nbclust(winedata1[,ncol(winedata1)], FUNcluster = kmeans, method = 'gap_stat')  
#From all of the cluster analysis the best are 2,3 and 4
```

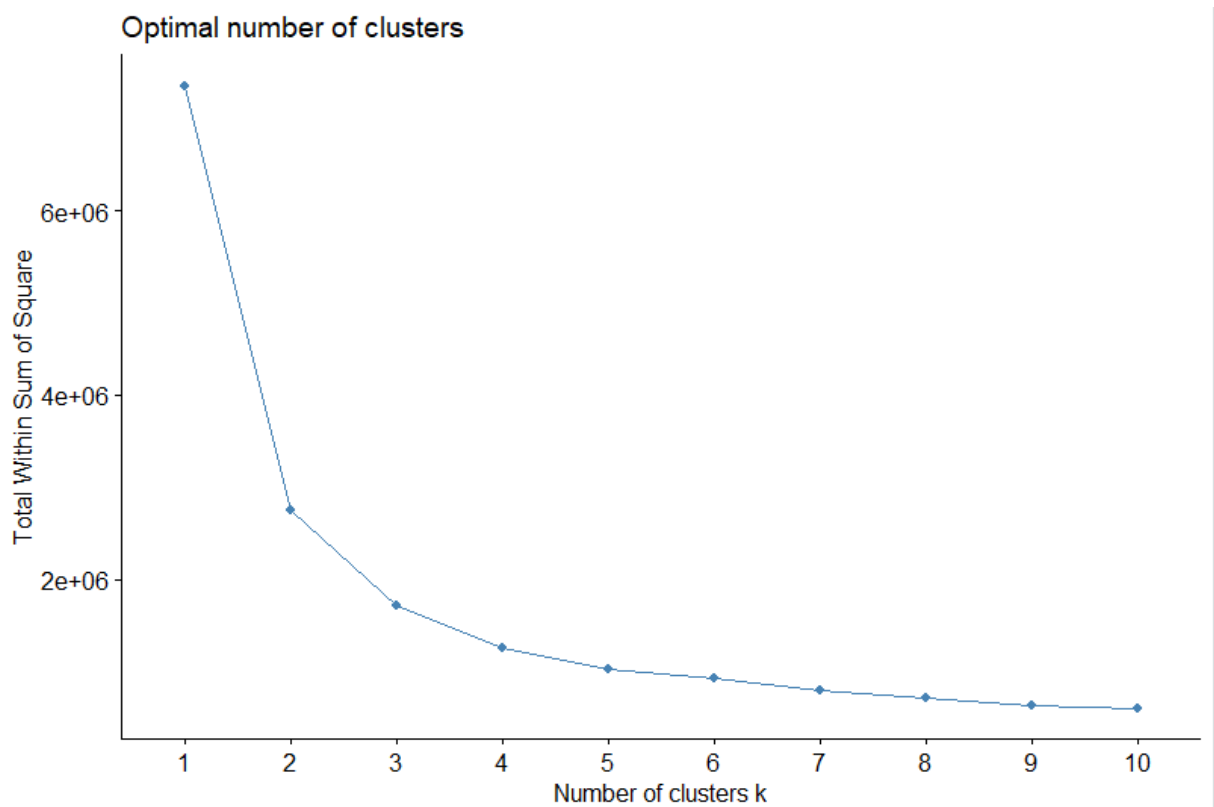


This graph shows that k=4 is the best possible cluster for the data





This graph shows that K=2 is the best choice in clustering



This graph shows that the optimal number of clusters would be k=2

Overall, from my findings I found out that 2,3 and 4 have all come up positively in all the manual and automatic tools.

### K-Means Analysis & final “winner” cluster case with providing brief explanation

#We now fit wine data to K-Means with k = 2

```
fit.km2 <- kmeans(winedata, 2)
```

#The original data set will be used to be compared with the 19th column

```
confuseTable.km <- table(winedata$quality, fit.km2$cluster)
```

confuseTable.km

	1	2
4	107	56
5	622	835
6	1271	927
7	639	241
8	120	55

The recall of sensitivity per class

R4       $107/163 = 65.64\%$

R5       $835/1457 = 57.30\%$

R6       $1271/2198 = 57.82\%$

R7       $639/880 = 72.61\%$

R8       $120/175 = 68.57\%$

Precision per cluster

Cluster 1  $2137/2759 = 77.45\%$

Cluster 2  $835/2114 = 39.49\%$

Accuracy of Clustering

$2972/4873 = 60.98\%$

```
#We now fit wine data to K-Means with k = 3

fit.km <- kmeans(winedata, 3)

#The original data set will be used to be compared with the 19th column
confuseTable.km <- table(winedata$quality, fit.km$cluster)

confuseTable.km
```

Here is the dominant winner for the 3 clusters

	1	2	3
4	91	36	36
5	368	570	519
6	836	887	475
7	412	394	74
8	80	75	20

The recall of sensitivity per class

R4       $91/163 = 55.82\%$

R5       $570/1457 = 39.12\%$

R6       $887/2198 = 40.35\%$

R7       $412/880 = 46.81\%$

R8       $80/175 = 45.71\%$

Precision per cluster

Cluster 1  $583/1787 = 32.62\%$

Cluster 2  $1457/1962 = 74.26\%$

Cluster 3  $0/1124 = 0\%$

Accuracy of Clustering

$2040/4873 = 41.86\%$

#We now fit wine data to K-Means with k = 4

```
fit.km3 <- kmeans(winedata, 4)
```

#The original data set will be used to be compared with the 19th column

```
confuseTable.km <- table(winedata$quality, fit.km3$cluster)
```

```
confuseTable.km
```

Here is the dominant winner for the 4 clusters

	1	2	3	4
4	41	27	64	31
5	392	350	202	513
6	786	328	450	634
7	398	51	219	212
8	81	13	39	42

The recall of sensitivity per class

R4       $64/163 = 39.26\%$

R5       $513/1457 = 35.20\%$

R6       $786/2198 = 35.75\%$

R7       $398/880 = 45.22\%$

R8       $81/175 = 46.28\%$

Precision per cluster

Cluster 1  $1247/1680 = 74.22\%$

Cluster 2  $0/769 = 0\%$

Cluster3  $64/974 = 6.5\%$

Cluster 4  $513/1432 = 35.82\%$

Accuracy of clustering

$1824/4855 = 37.56\%$

Overall, through all the calculations, I found out that cluster group 2 has the most accuracy out of all with cluster 1 being the winner of all the previous calculations. Against the 12<sup>th</sup> column the produced outputs which usually ended up on turning to 1% means that the data having little agreement between the wine quality and the cluster solution only being 0.01 and as calculated by the Kmeans analysis.

ARI

0.01126822

## Coordinates for each centre clustering group

Here we have the coordinates for the centre clustering of k=2 group

fit.km2\$centers #finds the Kmeans value and centralizes the data

```
> fit.km2$centers
  fixed.acidity volatile.acidity citric.acid residual.sugar
1    6.753218      0.2587645    0.3166881      4.866946
2    6.896292      0.2714504    0.3346292      8.491537
  chlorides free.sulfur.dioxide total.sulfur.dioxide
1 0.03935092      27.49550      108.5112
2 0.04648505      44.93331      176.8597
  density      pH sulphates      alcohol
1 0.9926785 3.188816 0.4718662 11.017582
2 0.9955204 3.184623 0.4943002  9.970752
```

Here we have the coordinates for the centre clustering of k=3 group

fit.km3\$centers #finds the Kmeans value and centralizes the data

```
  fixed.acidity volatile.acidity citric.acid residual.sugar
1    6.753645      0.2585117    0.3153980      4.108495
2    6.807803      0.2593376    0.3235541      6.786592
3    6.916205      0.2808209    0.3392324      9.323614
  chlorides free.sulfur.dioxide total.sulfur.dioxide
1 0.03814515      24.27525      96.34047
2 0.04332930      36.82197      141.52962
3 0.04733049      48.10075      194.47601
  density      pH sulphates      alcohol
1 0.9921804 3.188080 0.4714515 11.177124
2 0.9940914 3.188694 0.4758662 10.518533
3 0.9961731 3.182719 0.5058209  9.732605
```

Here we have the coordinates for the centre clustering of k=4 group

fit.km4\$centers #finds the Kmeans value and centralizes the data

```
  fixed.acidity volatile.acidity citric.acid residual.sugar
1    6.753799      0.2593307    0.3165123      5.482706
2    6.858243      0.2651404    0.3290399      7.803578
3    6.923053      0.2806794    0.3403206      9.566870
4    6.766241      0.2576450    0.3179930      3.577900
  chlorides free.sulfur.dioxide total.sulfur.dioxide
1 0.04046454      30.35022      121.18126
2 0.04555072      41.75634      158.36504
3 0.04747634      49.79924      203.11374
4 0.03728886      21.53016      85.03712
  density      pH sulphates      alcohol
1 0.9931284 3.195543 0.4753907 10.871322
2 0.9948924 3.189547 0.4824728 10.224801
3 0.9964087 3.176916 0.5095420  9.621196
4 0.9917996 3.178005 0.4675174 11.297970
```

## 2<sup>nd</sup> Objective (MLP)

The various methods that can be used for defining input vectors in time-series problems can be using an autoregressive approach as it's main purpose is related to being observing previous data and forecast any future data based on "a linear model, where current period values are a sum of past outcomes multiplied by a numeric factor" (Mehandzhiyski, V.,2021). This means that the autoregressive approach as its use isn't just limited one the day before but like a moving average it can go back to the last 7 days. There are also moving averages which can be used to calculate any amount of time periods specified as long as they are in an average for example a 7 day week would be considered a moving average and in time series problems it also assists in the time series data as smooths out any shorter term fluctuations in order to predict any future numerical data.

## References

- Zebari, R., Abdulazeez, A., Zeebaree, D., Zebari, D. and Saeed, J., 2020. A comprehensive review of dimensionality reduction techniques for feature selection and feature extraction. *Journal of Applied Science and Technology Trends*, 1(2), pp.56-70.
- Filter, L.V. and Filter, P.C.A., 2014. Seven Techniques for Dimensionality Reduction.
- KDnuggets. 2021. *Seven Techniques for Data Dimensionality Reduction - KDnuggets*. [online] Available at: <<https://www.kdnuggets.com/2015/05/7-methods-data-dimensionality-reduction.html>> [Accessed 15 November 2021].
- Karamizadeh, S., Abdullah, S.M., Manaf, A.A., Zamani, M. and Hooman, A., 2013. An overview of principal component analysis. *Journal of Signal and Information Processing*, 4(3B), p.173.
- Mehandzhiyski, V., 2021. What Is an Autoregressive Model? | 365 Data Science. [online] 365 Data Science. Available at: <<https://365datascience.com/tutorials/time-series-analysis-tutorials/autoregressive-model/>> [Accessed 10 November 2021].

## Appendix

### Part A

```
#Main r directory setup

setwd("~/Business intelligence cwk")

# Loading required packages

library(xlsx)

library(factoextra)

library(flexclust)

library(NbClust)

library(tidyr)

#Reading the winedata

winedata <- read.xlsx("Whitewine_v1.xlsx", sheetIndex = 1)
```

```
#Scaling and centralizing the data
```

```
winedata <- winedata
```

```
#Removes the 12th variable
```

```
winedata1 = winedata[,-12]
```

```
head(winedata1)
```

```
# checking for missing values
```

```
sum(is.na(winedata1))
```

```
summary(is.na(winedata1))
```

```
#Outlier detection by storing all numerical variables in an array structure
```

```
col = c("fixed.acidity", "volatile.acidity", "citric.acid", "residual.sugar", "chlorides", "free.sulfur.dioxide",  
"total.sulfur.dioxide", "density", "pH", "sulphates", "alcohol")
```

```
#Using the BoxPlot function to detect any presence of outliers in the data columns.
```

```
boxplot(winedata1[,c("fixed.acidity", "volatile.acidity", "citric.acid",  
"residual.sugar", "chlorides", "free.sulfur.dioxide", "total.sulfur.dioxide", "density", "pH", "sulphates",  
"alcohol")])
```

```
# The box plot identifies all present outliers
```

```
# All the outlier data from the numeric variables will now be replaced by NULL values
```

```
for (x in c("fixed.acidity", "volatile.acidity", "citric.acid",  
"residual.sugar", "chlorides", "free.sulfur.dioxide",  
"total.sulfur.dioxide", "density", "pH", "sulphates", "alcohol"))  
{  
  value = winedata1[,x][winedata1[,x] %in% boxplot.stats(winedata1[,x])$out]  
  winedata1[,x][winedata1[,x] %in% value] = NA  
}
```

```
#Checking whether the outliers have been replaced by null data or not
```

```
as.data.frame(colSums(is.na(winedata1)))
```

```
#This is responsible for removing all the null values
```

```
winedata1 = drop_na(winedata1)
```

```
as.data.frame(colSums(is.na(winedata1)))
```

```
#Output a newly created boxplot without outliers
```

```
boxplot(winedata1)
```

```
#Manually finding clusters 1
```

```
wss <- 0
```

```
for (i in 1:12){
```

```
  wss[i] <-
```

```
    sum(kmeans(winedata1, centers=i)$withinss)
```

```
}
```

```
plot(1:12,
```

```
  wss,
```

```
  type="b", ### "b" for both####
```

```
  xlab="Number of Clusters",
```

```
  ylab="Within groups sum of squares")
```

```
#Manually finding clusters in a larger criteria
```

```
set.seed(1234)
```

```
nc <- NbClust(winedata1,
```

```
  min.nc=2, max.nc=15,
```

```
  method="kmeans")
```

```
table(nc$Best.n[1,])
```



```

barplot(table(nc$Best.n[1,]), # provide bar charts####
        xlab="Numer of Clusters",
        ylab="Number of Criteria",
        main="Number of Clusters Chosen by 30 Criteria")

```

```

# Automatically finding clusters from 3 different methods
fviz_nbclust(winedata1[, -ncol(winedata1)], FUNcluster = kmeans, method = 'wss')
fviz_nbclust(winedata1[, -ncol(winedata1)], FUNcluster = kmeans, method = 'silhouette')
fviz_nbclust(winedata1[, -ncol(winedata1)], FUNcluster = kmeans, method = 'gap_stat')
#From all of the cluster analysis the best are 2,3 and 4

```

```

#We now fit wine data to K-Means with k = 3
fit.km <- kmeans(winedata, 3)
#The original data set will be used to be compared with the 19th column
confuseTable.km <- table(winedata$quality, fit.km$cluster)
confuseTable.km
#Checking specific information relating to the centre clustering
fit.km$centers

```

```

#We now fit wine data to K-Means with k = 2
fit.km2 <- kmeans(winedata, 2)
#The original data set will be used to be compared with the 19th column
confuseTable.km <- table(winedata$quality, fit.km2$cluster)
confuseTable.km
#Checking specific information relating to the centre clustering
fit.km2$centers

```

```

#We now fit wine data to K-Means with k = 4
fit.km3 <- kmeans(winedata, 4)

```

```
#The original data set will be used to be compared with the 19th column  
confuseTable.km <- table(winedata$quality, fit.km3$cluster)  
confuseTable.km  
#Checking specific information relating to the centre clustering  
fit.km3$centers
```

Part B